

RouterEval real LLM pools, strong_to_weak (10-1,000) + leaderboard (5,000), honest ($\beta = 0$)

