


RouterEval real LLM pools, all_weak (10-1,000) + leaderboard (5,000), $\beta = 0.5$, random liars



■ warm-start bandit
 ■ MIDIAN-A
 ■ MIDIAN
 ■ flat probe argmax (online)
 ■ flat probe argmax frozen
 ■ RouterBench KNN router
 ■ linucb honest
 ■ random
 ■ knn router
 ■ MIDIAN-V