


RouterEval real LLM pools, strong_to_weak (10-1,000) + leaderboard (5,000), $\beta = 0.5$, random liars



RouterBench MLP router, declared argmax, MIDIAN-V, MIDIAN, flat probe argmax frozen, linucb honest, autogen, llama, langgraph, random, cr, MIDIAN-VA, warm-start bandit, MIDIAN-A, flat probe argmax (online), RouterBench KNN router, adk, camel, maf, openai, smol, kn