


RouterEval real LLM pools, strong_to_weak (10-1,000) + leaderboard (5,000), $\beta = 0.5$, random liars



le
halving (peer)

RouterBench MLP router
MIDIAN-V
MIDIAN-A
MIDIAN
flat probe argmax (online)
midian sh
declared argmax
warm-start bandit
flat probe argmax frozen
RouterBench KNN router
linucb honest
knn router
random