


RouterEval real LLM pools, strong_to_weak (10-1,000) + leaderboard (5,000), $\beta = 0.5$, random liars



(peer)

RouterBench MLP router

declared argmax

MIDIAN-V

MIDIAN

flat probe argmax frozen

RouterBench KNN router

linucb honest

autogen

llama

langgraph

random

cr

MIDIAN-VA

warm-start bandit

MIDIAN-A

flat probe argmax (online)

adk

camel

maf

openai

smol

kn